

# P4: Portable Parallel Processing Pipelines for Interactive Information Visualization

**Kelvin Li**

University of California, Davis

BayVAST, 2018

# Goals

- Combine declarative visualization design and GPU Computing
- Provide both performance and flexibility
- Support developing high-performance visualization systems for big data applications

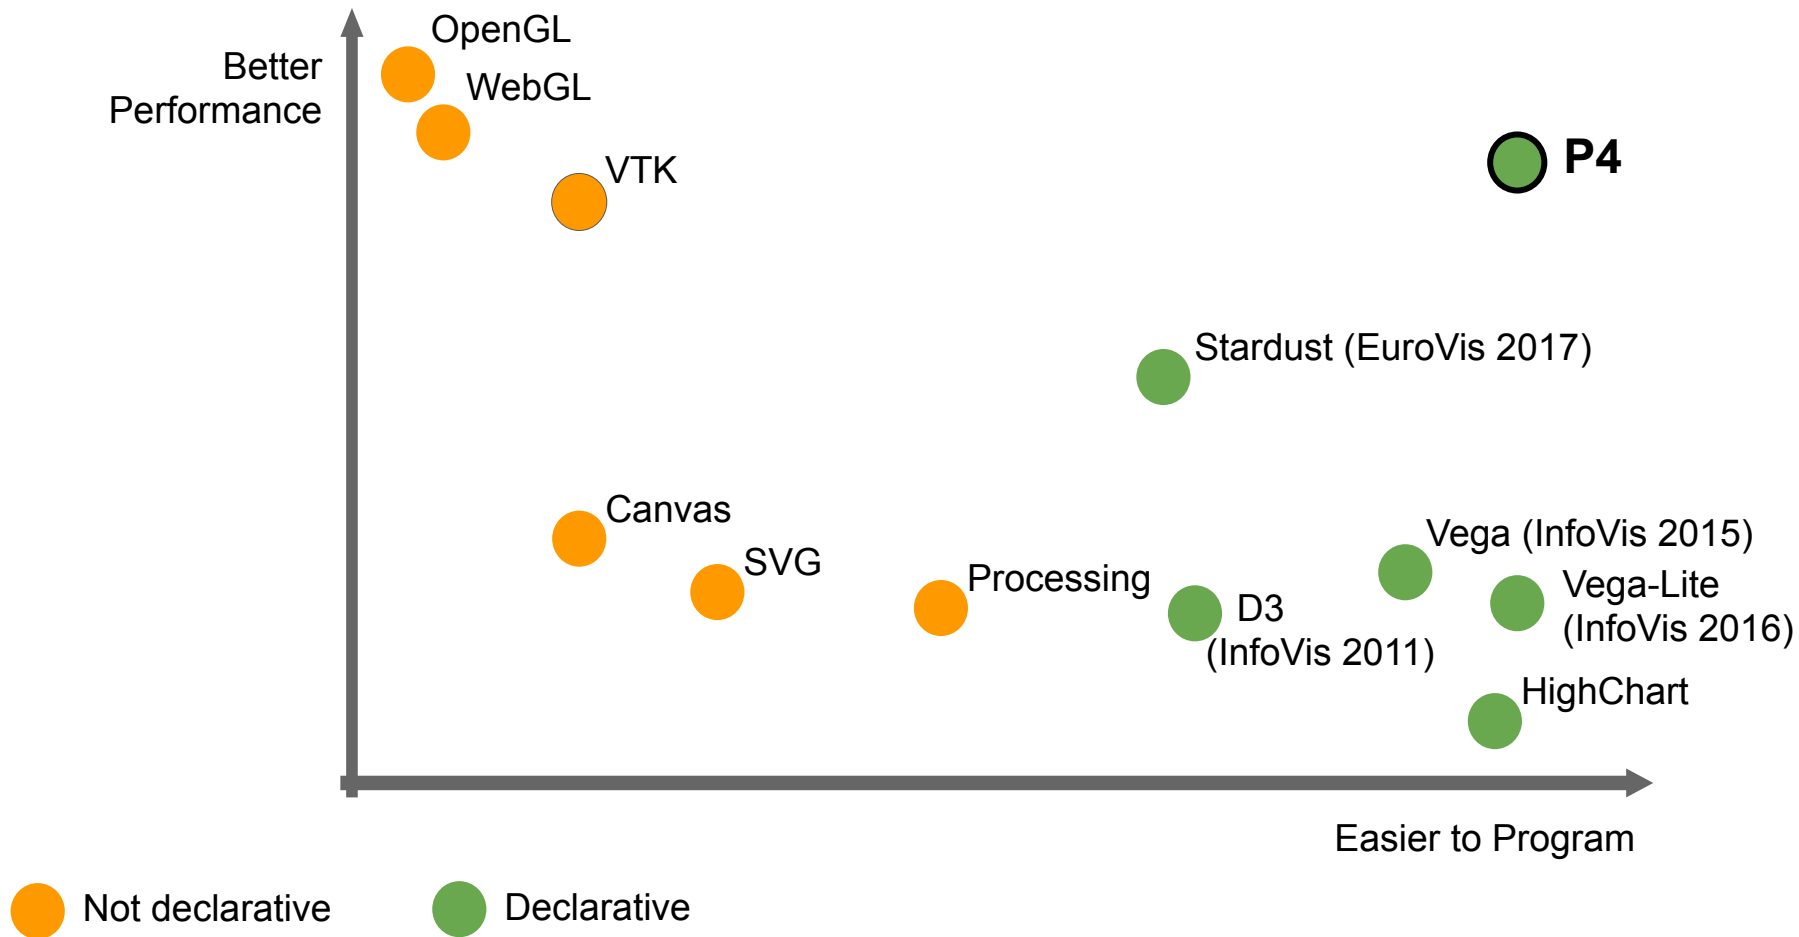
# Background: Declarative Visualization Language

Declarative visualization languages specify ***what*** should be visualized instead of ***how*** to visualize.

- Make visualization systems much easier to develop
- Become popular for creating visualization applications
  - Examples: ggplot and D3

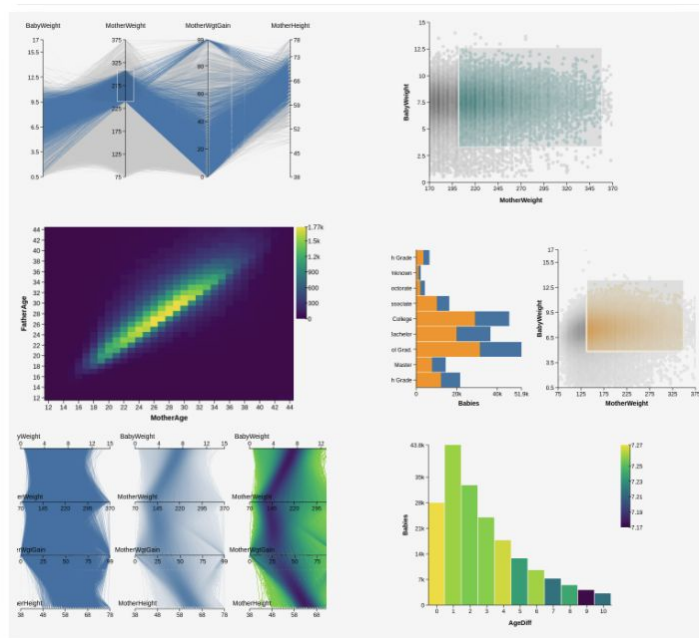
# Visualization Toolkits and Libraries

Declarative visualization languages (e.g., D3, ggplot, Vega/Vega-Lite)	Easy to use but Low performance
Low-level graphics and visualization libraries (e.g., OpenGL/WebGL, VTK)	High performance but Difficult to program

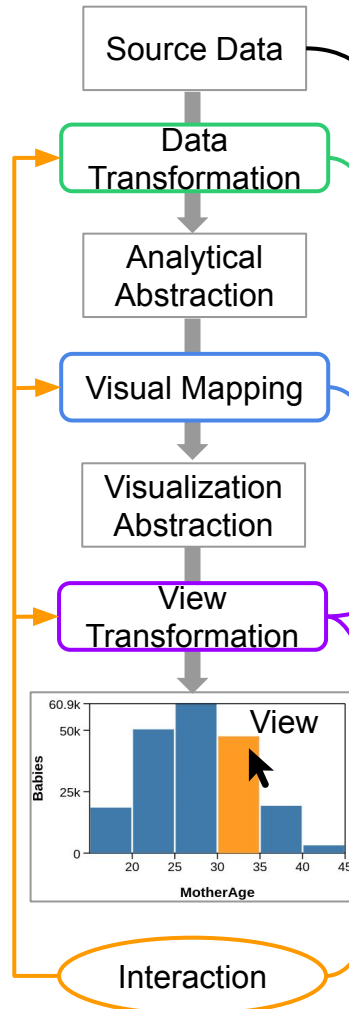


# Current Version of P4


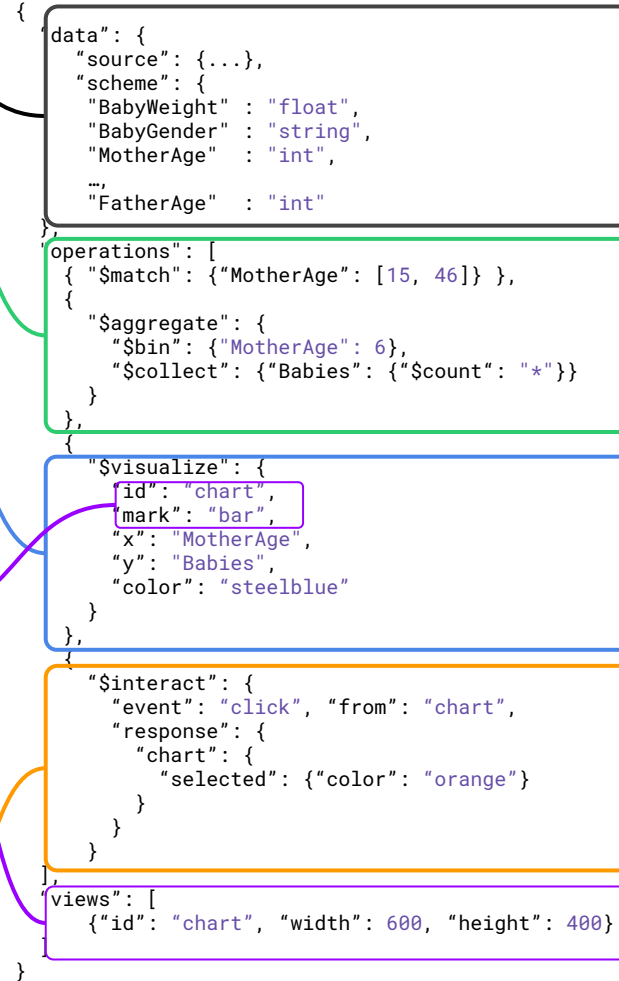
- Implemented based on WebGL 1.0
- Provided most common data transformations (e.g., *deriving new values, filtering, binning, and group-by aggregations*)
- Supported common visualization designs (e.g., *bar charts, heat maps, scatter plots, parallel coordinates*)



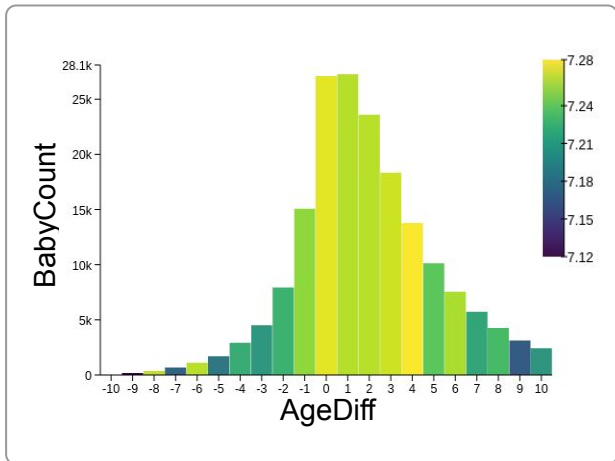
## InfoVis Reference Model



## P4 Declarative Grammar



# JavaScript API

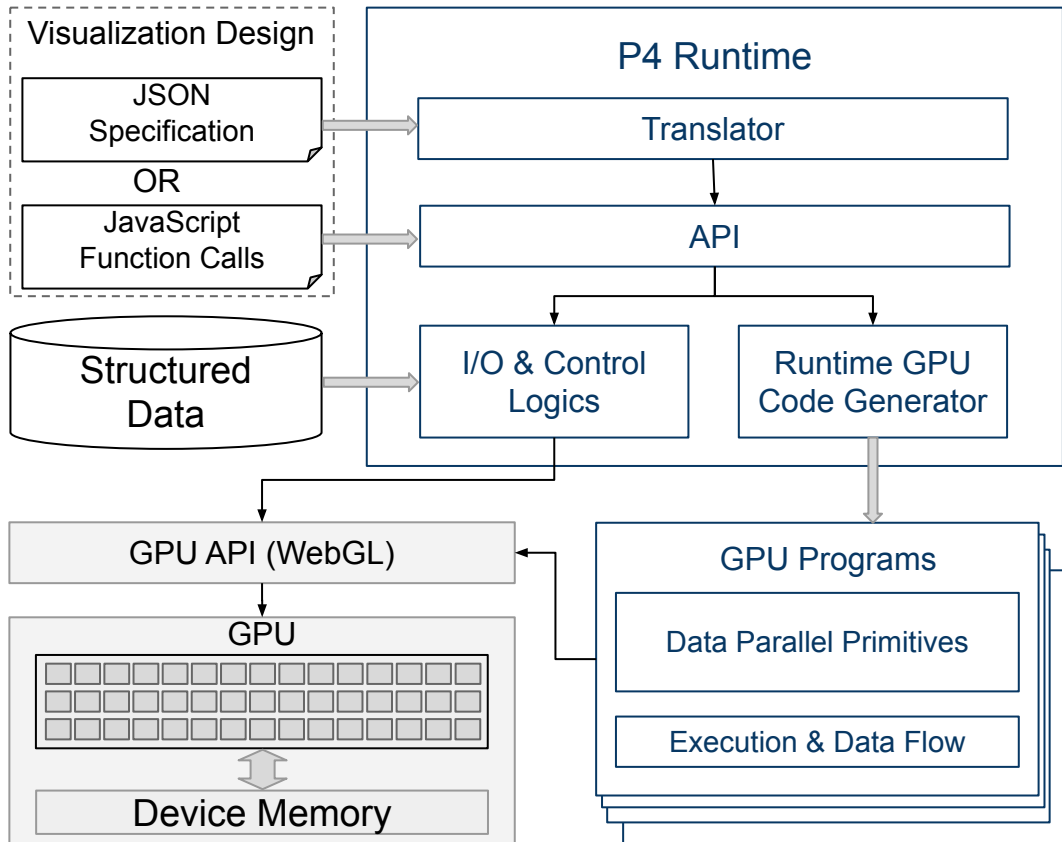


```
pipeline.data(...)
  .derive({ AgeDiff: "FatherAge - MotherAge" })
  .match({ AgeDiff: [-10, 10] })
  .aggregate({
    $group: "AgeDiff",
    $collect: {
      BabyCount: { $count: "*" },
      AvgBabyWeight: { $avg: "BabyWeight" }
    }
  })
  .visualize({
    mark: "bar",
    x: "AgeDiff",
    y: "BabyCount",
    color: {
      field: "AvgBabyWeight",
      scheme: "viridis"
    }
  })
})
```



# Dataflow and Workflow

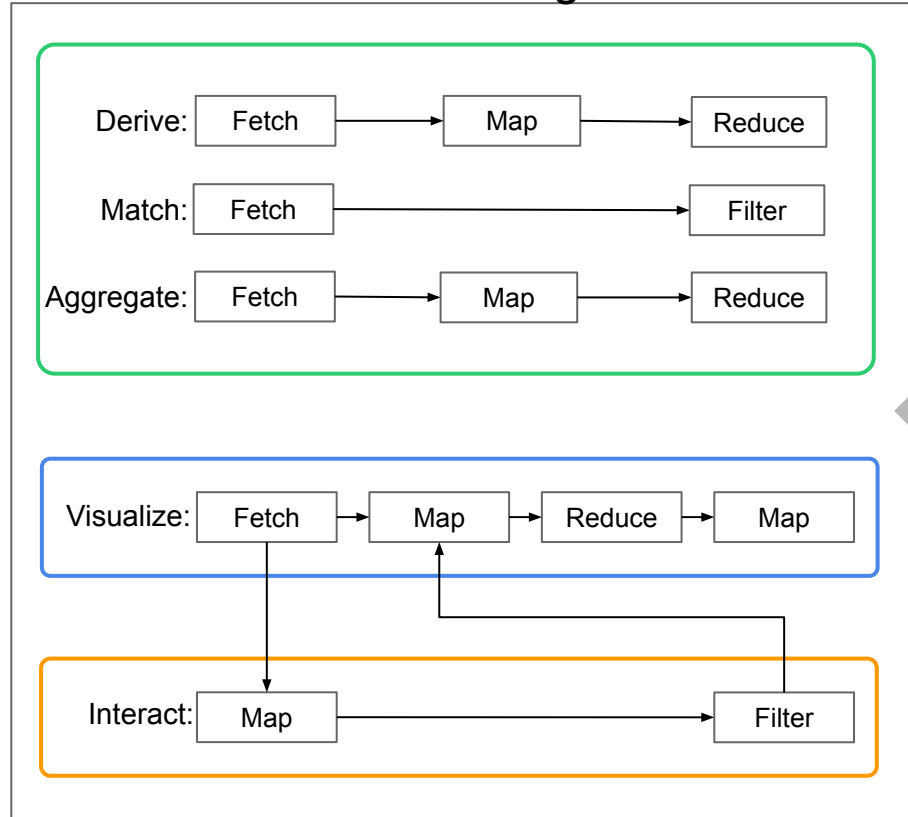
- Cache all data in GPU memory
- Accelerate both data transformation and visualization operations using the GPU



# Runtime GPU Code Generator


- Create GPU programs based on user specifications at runtime
- Compose all data processing and visualization operations by using four *data-parallel primitives (Fetch, Map, Filter, and Reduce)*

## P4 Parallel Processing Framework




## Unified Data Model


### Packaged Data in Texture



### Metadata



### Filter Result



## Aggregation Query

```
var aggr = P4.pipeline().data(...)  
  .aggregate({  
    $group: "MotherRace",  
    $collect: {  
      AvgWeight: {  
        $avg: "BabyWeight"  
      }  
    }  
  })
```

## WebGL Pipeline

### GPU Memory

### Textures

## Generated WebGL 1.0 Code

```
varying float value;  
uniform int p4_aggr_key;  
uniform int p4_field_id;  
...  
float p4_dpp_fetch(int key) {...};  
float p4_dpp_map(int key) {...};
```

c1

```
P4.pipeline().data(...)  
  .aggregate({  
    $group: "MotherRace",  
    $collect: {  
      AvgWeight: {  
        $avg: "BabyWeight"  
      }  
    }  
  })
```

```
t(p4_dpp_fetch(p4_aggr_key));  
p_fetch(p4_field_id);  
_dpp_map(p4_aggr_key);  
vec4(ox*2.-1., 0., 0., 1.);
```

c2

```
or = vec4(0., 0., 1., value);
```

c3

```
var gl = aggr.ctx;  
gl.enable(gl.BLEND);  
gl.blendFunc(gl.ONE, gl.ONE);  
gl.blendEquation(reduce_method); //gl.ADD
```

c4

```
gl.drawArraysInstanced(gl.POINTS, ...,
```


## Aggregation Workflow

### Texture



### Fetch

### Vertex Space




### Map

### Pixel Space



### Reduce

### Frame Buffer




## Blender

## Offscreen FBO



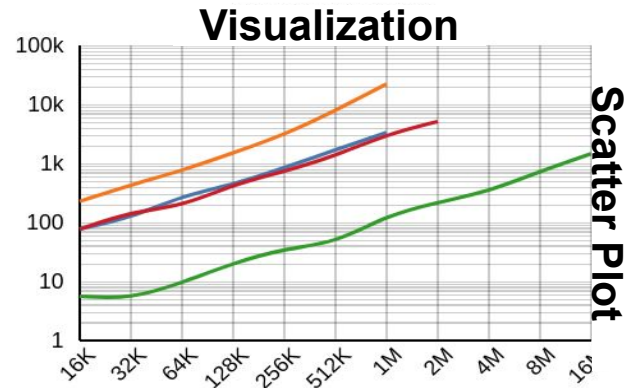
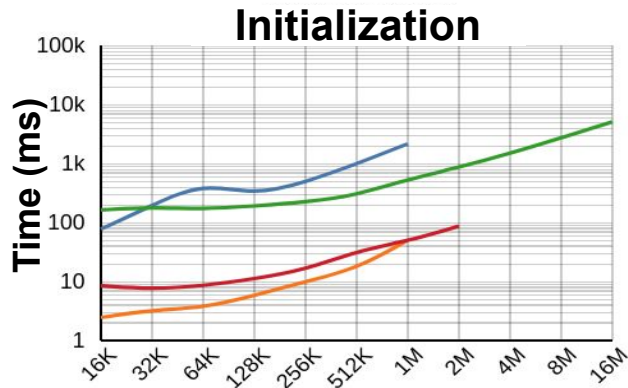
## Views

# Benchmarks: Data Transformation



Benchmarks performed using Google Chrome on a system with a Intel i7-4790 CPU and Nvidia GTX Titan GPU

# Benchmarks: Visualization



Number of Data Records

# Limitation and Future Work

- Data size must be less than GPU memory limit
- Current version only supports a small set of data transformation and visualization operations
- Future work:
  - Support progressive visual analytics to allow data size larger than memory
  - Add more functionalities for supporting visual analytics
  - Allow extensions to be added by others

# Conclusion

- P4 combines declarative visualization design and GPU computing
- P4's novel framework makes this possible:
  - Formulate all InfoVis operations using only four data-parallel primitives (DPP)
  - Leverage **DPP** to generate GPU code based on user specification at runtime
- Our framework can be extended and applied to:
  - Server backends and desktop environments
  - Distributed computing on GPU clusters



# Thank You !

Email: [kelli@ucdavis.edu](mailto:kelli@ucdavis.edu)

Project Homepage: <https://jpkli.github.io/p4>

Source Code: <https://github.com/jpkli/p4>

TVCG Paper: <https://ieeexplore.ieee.org/document/8468065>